

UNIVERSITY OF RWANDA

COLLEGE OF SCIENCE AND TECHNOLOGY
Y3 CSE



Group 4: TransConnect

Members: MUSAZA PATRICK (223019061)
MANZI NSENGA IVAN (223004392)

Module: Computing Intelligence & Applications
Date: 11th December 2025

Introduction

TransConnect is a smart AI web-based transit prediction platform designed to provide:

-  Predicted-time bus movement tracking
-  Predictive travel times using machine learning
-  More accurate ETA (Estimated Time of Arrival) predictions
-  The system improves commuter experience by reducing uncertainty and wait time



Project Objectives

-  Provide accurate travel time predictions for public transport
-  Reduce commuter waiting time and uncertainty
-  Offer traffic-adjusted ETA during rush hours
-  Improve the efficiency of the (Kabuga–Nyabugogo route) case study
-  Use machine learning to support data-driven transport planning

Background

Public transport users
in Rwanda often
experience:

- Unpredictable delays
 - Lack of predicted-time transit information
 - Traffic-based inconsistencies, especially in mornings and evenings
- The Kabuga–Nyabugogo case study route is a major corridor, making it ideal for predictive analytics

Rationale

Why TransConnect?



Passengers need accurate ETA to plan their journeys



Transport operators need better route visibility



Machine Learning can reduce uncertainty and enhance mobility



Real-time predictions = better satisfaction + optimized transport flow





System Features

- Real-time GPS updates (to be implemented later)
- Distance-to-destination calculation
- ML-based ETA prediction (Random Forest, Gradient Boosting, Regression & LSTM).
- Traffic condition classification (normal, rush hour)
- Remaining distance + arrival time display

Machine Learning Models

- Models evaluated:
 - Linear Regression
 - Decision Tree
 - Random Forest
 - Gradient Boosting
 - Long Short-Term Memory
- Best model is automatically selected based on highest R^2 and lowest RMSE



Expected Impact

- Improved commuter planning
- Reduced delays and congestion
- Increased operational transparency
- Smart city readiness for public transport



Conclusion

TransConnect demonstrates how Machine Learning can transform public transportation by:

- Providing real-time predictive insights
- Enhancing user satisfaction
- Supporting data-driven decision making
- Contributing to Rwanda's smart mobility vision

DATASET OVERVIEW

Source: Bus route information for Kigali public transport



Data Description:

- Tabular data with bus routes, stop names, and coordinates
- Multiple routes with sequential stops
- Mixed data types: text (stop names) and numeric (coordinates)

Initial Data Quality Issues:

- Inconsistent stop naming conventions
- Mixed coordinate formats S139282 E30.28382)
- Potential missing values and duplicates
- Non-standardized data structure across routes
- Nature: Spatial-Temporal Data with Geographic Coordinates



H16

fx

DATA PREPOCESSING STEPS

1. DATA CLEANING

Purpose: Handle missing, incorrect, and inconsistent data

Application to TransConnect:

- Identify missing values
- Clean coordinate formats: Remove 'E', 'S' characters
- Convert to standard decimal degrees: -
1.39282, 30.28382
- Handle missing stop names or coordinates
- Handle duplicates

```
# Load the route details CSV file
df = pd.read_csv('Route Details - Kabuga - Nyabugogo.csv')

# Clean coordinate formats - convert S/E to negative/positive
df_clean['latitude'] = df_clean['Latitude'].str.replace('S', '-').astype(float)
df_clean['longitude'] = df_clean['Longitude'].str.replace('E', '').astype(float)

# Handle missing values
df_clean = df_clean.dropna(subset=['latitude', 'longitude', 'Stops'])

# Remove duplicates based on stop name and coordinates
df_clean = df_clean.drop_duplicates(subset=['Stops', 'latitude', 'longitude'])
```

2. DATA INTEGRATION

```
# Create unique stop IDs
df_clean['stop_id'] = range(1, len(df_clean) + 1)

# Assign route information
df_clean['route_id'] = 'Kabuga_Nyabugogo_001'
df_clean['route_name'] = 'Kabuga - Nyabugogo'

# Create route-stop sequence mapping
route_stop_sequence = {
    'Kabuga_Nyabugogo_001': df_clean[['stop_id', 'stop_sequence',
'Stops']].sort_values('stop_sequence').to_dict('records')
}
```

Data Integration:

- Merge multiple route datasets into unified structure
- Create stops table (unique stops with IDs)
- Create routes table (route sequences referencing stops)

Impact:

- Efficient multi-route management

3. DATA REDUCTION

```
# Select key features for modeling
essential_features = [
    'route_id', 'route_name', 'stop_id', 'stop_sequence',
    'latitude', 'longitude', 'Stops'
]

df_reduced = df_integrated[essential_features].copy()

# Rename columns for consistency
df_reduced = df_reduced.rename(columns={'Stops': 'stop_name'})

# Remove irrelevant columns that were in raw data
columns_removed = set(df_integrated.columns) -
set(essential_features)
```

Data Reduction:

- ❑ Remove irrelevant columns: "Route Price"
- ❑ Focus on core features: stop_id, latitude, longitude
- ❑ Deduplicate identical stops across routes

Impact:

- Faster processing for route optimization
- Foundation for transfer point identification

4. DATA TRANSFORMATION

Purpose: Convert data into model-ready format



Key Transformations for TransConnect:

- Coordinate validation and normalization
- Feature engineering using Haversine formula:
 - Distance between consecutive stops
 - Estimated travel time between stops
 - Create sequential ordering of stops



Impact:

Raw coordinates actionable features for AI model

```
R = 6371 # Earth radius in km

# Convert degrees to radians
lat1, lon1, lat2, lon2 = map(radians, [lat1, lon1,
lat2, lon2])
dlat = lat2 - lat1
dlon = lon2 - lon1

# Haversine formula
a = sin(dlat/2)**2 + cos(lat1) * cos(lat2) *
sin(dlon/2)**2
c = 2 * atan2(sqrt(a), sqrt(1-a))
```

```
# Update distance to next stop
df_reduced.loc[current_idx,
'distance_to_next'] = distance

# Estimate travel time (assuming
average speed 25 km/h in urban areas)
travel_time = (distance / 25) * 60 # Convert to minutes
df_reduced.loc[current_idx,
'estimated_travel_time_min'] = travel_time
```

HAVERSINE FORMULA

$$d = 2r \arcsin \left(\sqrt{\sin^2 \left(\frac{\phi_2 - \phi_1}{2} \right) + \cos(\phi_1) \cos(\phi_2) \sin^2 \left(\frac{\lambda_2 - \lambda_1}{2} \right)} \right)$$

5. DATA DISCRETIZATION

```
# Discretize cumulative distance into route segments
max_distance = df_transformed['cumulative_distance'].max()
distance_bins = [0, max_distance*0.25, max_distance*0.5,
max_distance*0.75, max_distance + 0.1]
distance_labels = ['Start', 'Early', 'Mid', 'Late']

df_transformed['route_segment'] = pd.cut(
    df_transformed['cumulative_distance'],
    bins=distance_bins,
    labels=distance_labels,
    include_lowest=True
)
```

```
# Discretize travel time estimates - handle NaN values
time_bins = [0, 2, 5, 10, float('inf')]
time_labels = ['Very Short', 'Short', 'Medium', 'Long']

df_transformed['travel_time_category'] = pd.cut(
    df_transformed['estimated_travel_time_min'],
    bins=time_bins,
    labels=time_labels
)
```

Data Discretization:

Purpose: Convert continuous data into categories for simpler analysis.

Convert continuous "distance_from_start" into categories:

- Bin 1: "Start" (0–2 km)
- Bin 2: "Early" (2–5 km)
- Bin 3: "Mid" (5–10 km)
- Bin 4: "Late" (10+ km)

Enables traffic pattern analysis by route segment

Impact: Helps analyze congestion patterns or route efficiency per time segment.

6. DATA AUGMENTATION

```
# Add condition column to original  
data first  
df_processed = df_processed.copy()  
df_processed['condition'] = 'normal'  
augmented_data = [df_processed] #  
Start with original data
```

Data Augmentation:


- Generate synthetic trips with varying travel times
- Simulate different conditions: rush hour, weather, events
- Add temporal features: time_of_day, day_of_week

Purpose: Enrich the dataset with synthetic variations to improve prediction accuracy.

Impact:

More robust and reliable system, even under uncertain or missing data conditions.

5. VISUALIZATION & IMPACT



Before Preprocessing:

- Messy, overlapping points on map
- Inconsistent stop representations
- Unreliable spatial relationships



After Preprocessing:

- Clean, sequential route visualization
- Accurate stop positioning
- Calculated distances and travel times



Impact on Model Readiness:

- Clean spatial data for route optimization
- Engineered features for arrival prediction
- Structured data for CI algorithms (Genetic Algorithms, ML)
- Foundation for real-time intelligence system

BEFORE – DATA PREPROCESSING

	A	B	C	D
1	Stops	Latitude	Longitude	
2	Kabuga Bus Park1	S1.97922	E30.22352	
3	Masaka Hospital	S1.98699	E30.21955	
4	Masaka Bus Terminal	S1.99804	E30.19164	
5	Masaka Rd_2	S1.99587	E30.19155	
6	Masaka Health Centre	S1.99302	E30.19111	
7	Masaka Rd_1	S1.99001	E30.19285	
8	At 19_Inyange1	S1.98532	E30.19108	
9	KK 3 Rd_15	S1.97920	E30.18248	
10	Kibaya_Minagri	S1.97749	E30.18150	
11	Bus Route 107_3	S1.97365	E30.17860	
12	Mulindi1	S1.97129	E30.17572	
13	KK 3 Rd_11	S1.96891	E30.16975	
14	At 15 Ndera1	S1.96715	E30.16319	
15	SEZ B2	S1.95801	E30.15010	
16	Kigali Parents School2	S1.95409	E30.13899	
17	La Parisse1	S1.95748	E30.13211	
18	Ku Cyamutzig	S1.96272	E30.12460	
19	Prince House2	S1.96011	E30.11533	
20	Alpha Palace Hotel	S1.96120	E30.11216	
21	Good Year2	S1.96281	E30.10898	
22	Sonatubes4	S1.96820	E30.10170	
23	Bralirwa3	S1.96891	E30.09799	
24	Amasezerano	S1.96807	E30.09426	
25	Rwandex3	S1.96918	E30.08676	
26	Kwa Mironko2	S1.96752	E30.08285	
27	KN 3 Rd_6	S1.96489	E30.08039	
28	Volta Super1	S1.96249	E30.07761	
29	Ku Mazi1	S1.95730	E30.07608	
30	Chez Rasta	S1.95026	E30.07484	
31	One Love	S1.94667	E30.07381	
32	Forex Form_Poids Lourd	S1.94101	E30.06934	
33	Kigali Gaz_Kinamba	S1.93686	E30.06433	

DATA PREPOCESSING ACTIONS

== DATA ASSESSMENT ==

Dataset shape: (36, 3)

Columns: ['Stops', 'Latitude', 'Longitude']

First 5 rows:

	Stops	Latitude	Longitude
0	Kabuga Bus Park1	S1.97922	E30.22352
1	Masaka Hospital	S1.98699	E30.21955
2	Masaka Bus Terminal	S1.99804	E30.19164
3	Masaka Rd_2	S1.99587	E30.19155
4	Masaka Health Centre	S1.99302	E30.19111

Missing values:

Stops 0
Latitude 0
Longitude 0
dtype: int64

== DATA CLEANING ==

Original stops: 36

Cleaned stops: 36

Sample cleaned data:

	Stops	latitude	longitude
0	Kabuga Bus Park1	-1.97922	30.22352
1	Masaka Hospital	-1.98699	30.21955
2	Masaka Bus Terminal	-1.99804	30.19164
3	Masaka Rd_2	-1.99587	30.19155
4	Masaka Health Centre	-1.99302	30.19111

== DATA INTEGRATION ==

Integrated dataset shape: (36, 9)

Routes processed: 1

Stops in route: 36

== DATA REDUCTION ==

Columns removed: {'Longitude', 'Latitude'}

Reduced dataset shape: (36, 7)

Final features: ['route_id', 'route_name', 'stop_id', 'stop_sequence', 'latitude', 'longitude', 'stop_name']

== FEATURE ENGINEERING ==

Processing route: Kabuga_Nyabugogo_001 with 36 stops

Feature engineering completed:

- Total route distance: 26.07 km
- Total estimated time: 62.6 minutes
- Average distance between stops: 0.72 km
- Average time between stops: 1.7 min

== DATA DISCRETIZATION ==

Discretization completed:

Route segments distribution:

route_segment

Start	count
Early	8
Mid	9
Late	12

Name: count, dtype: int64

Travel time categories:

travel_time_category

Category	count
Very Short	26
Short	9
Medium	1
Long	0

Name: count, dtype: int64

Route segments distribution:

route_segment

Category	count
Start	42
Early	48
Mid	59
Late	67

Name: count, dtype: int64






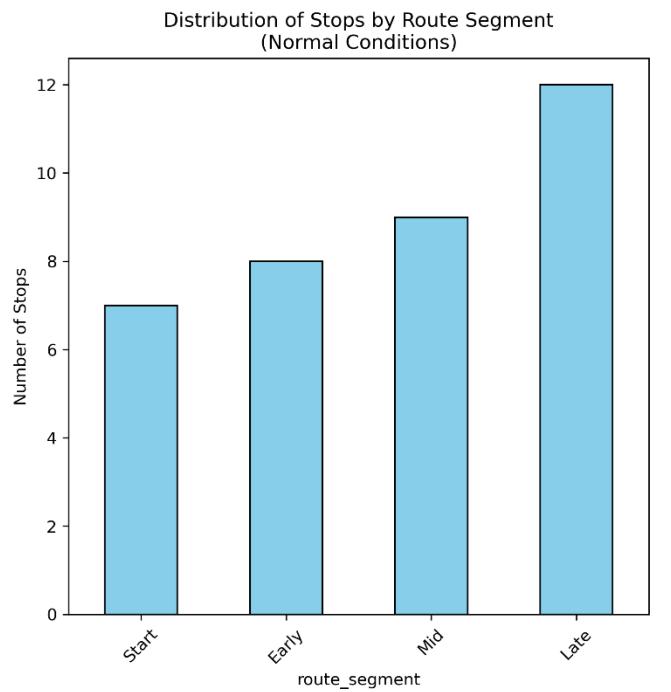
Travel time categories:

travel_time_category

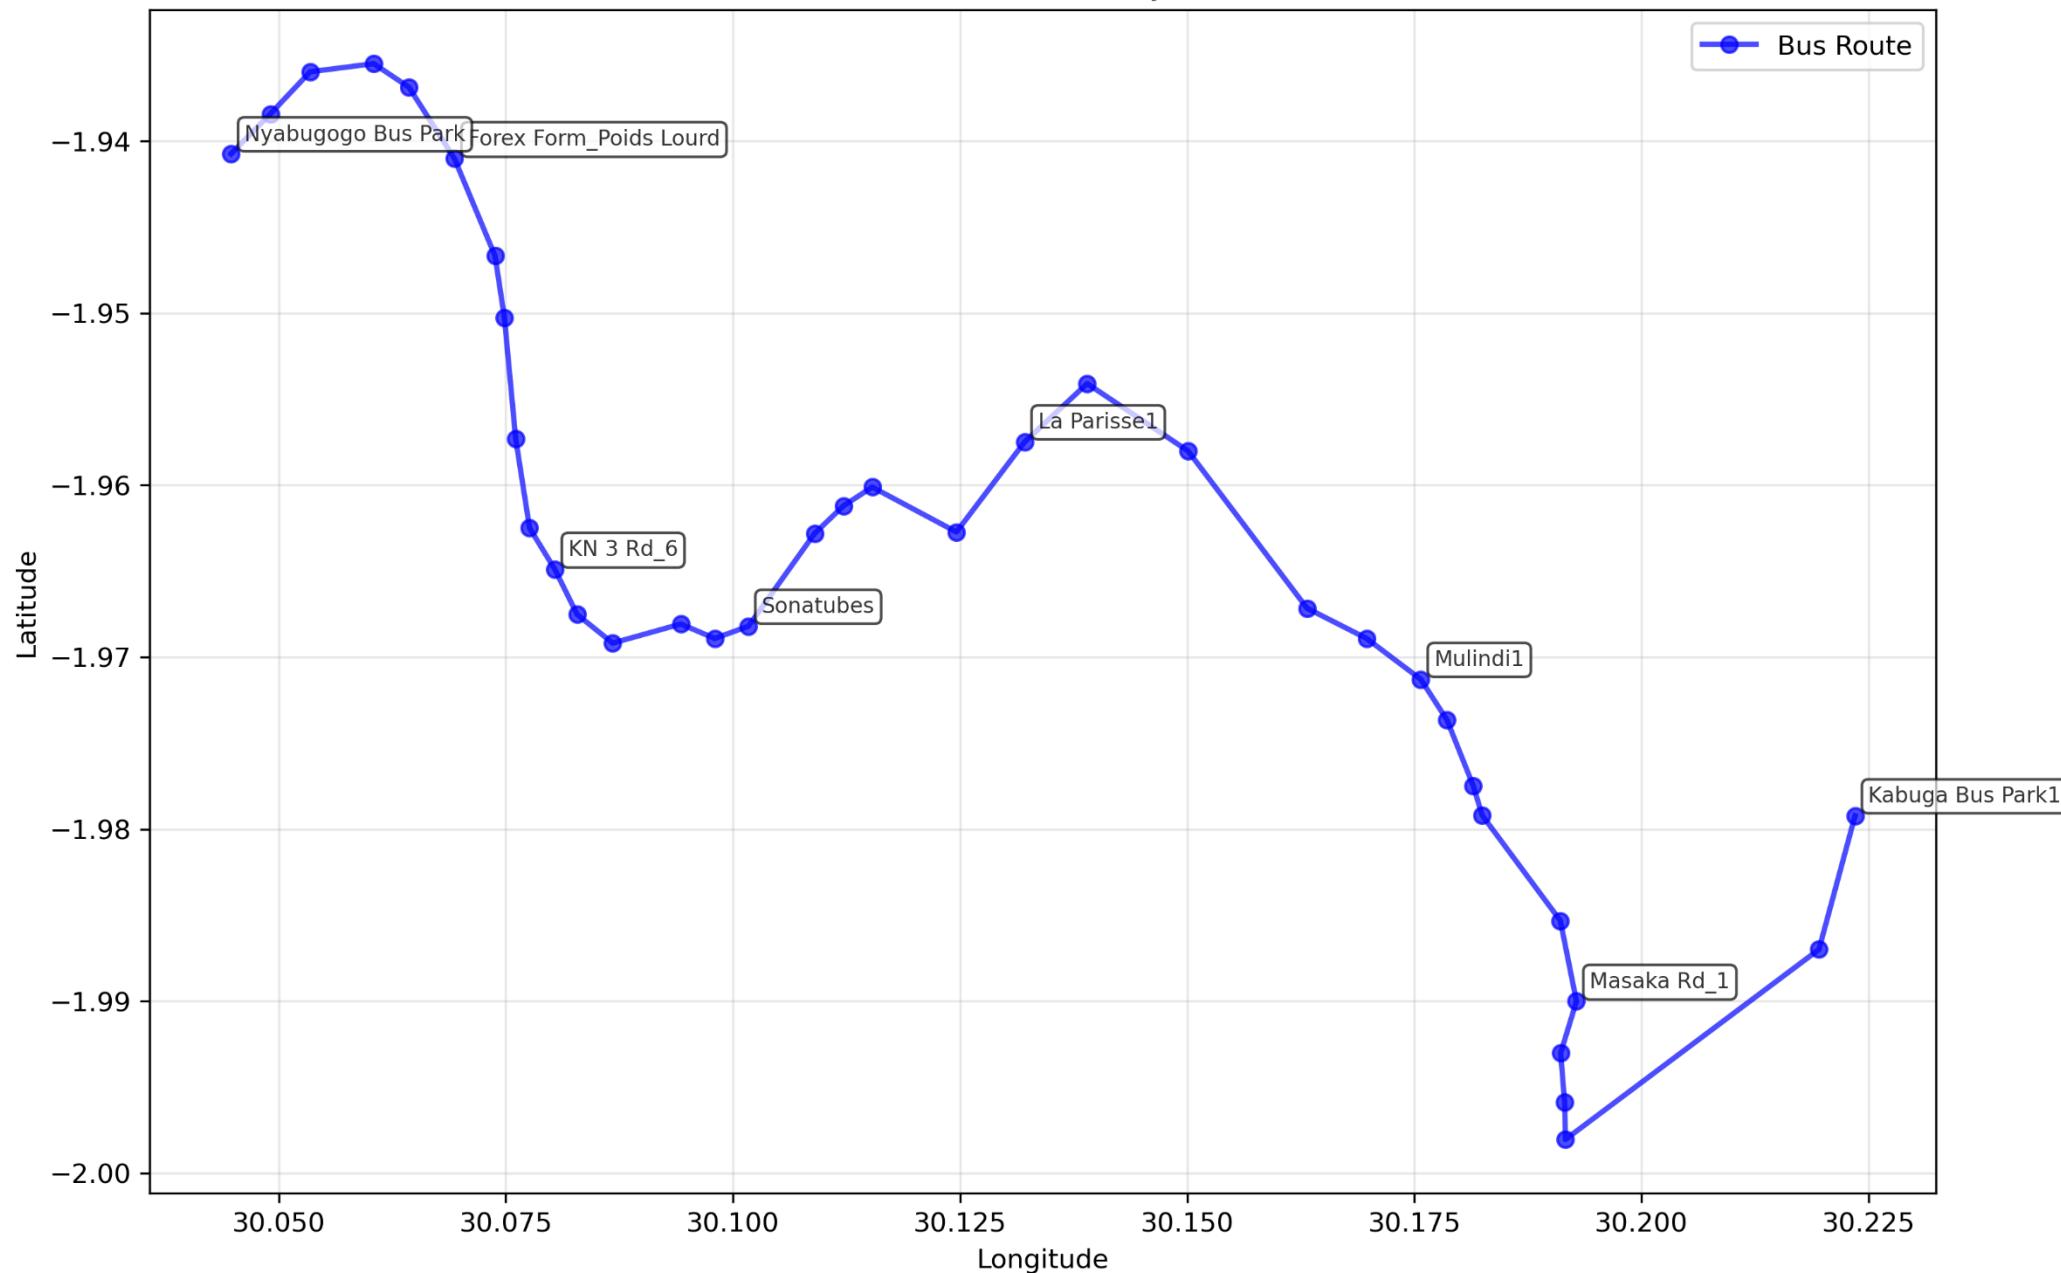
Category	count
Very Short	135
Short	68
Medium	10
Long	3

Name: count, dtype: int64

VISUALIZATION



Kabuga - Nyabugogo Bus Route TransConnect Project



AFTER – CLEAN PREPROCESSING

	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1	route_id	route_name	stop_id	stop_sequence	latitude	longitude	stop_name	distance_to_start	cumulative_distance	estimated_travel_time	travel_time_category	travel_time_cat	distance_category	condition
2	Kabuga	Kabuga - Nyabu	1	1	-1.979	30.224	Kabuga Bus P	0.970107559	0	2.328258142	Start	Short	Close	normal
3	Kabuga	Kabuga - Nyabu	2	2	-1.987	30.22	Masaka Hospital	3.33608671	0.970107559	8.006608105	Start	Medium	Far	normal
4	Kabuga	Kabuga - Nyabu	3	3	-1.998	30.192	Masaka Bus Ter	0.24150018	4.306194269	0.579600431	Start	Very Short	Very Close	normal
5	Kabuga	Kabuga - Nyabu	4	4	-1.996	30.192	Masaka Rd_2	0.320655506	4.547694449	0.769573213	Start	Very Short	Very Close	normal
6	Kabuga	Kabuga - Nyabu	5	5	-1.993	30.191	Masaka Health	0.386537039	4.868349955	0.927688894	Start	Very Short	Very Close	normal
7	Kabuga	Kabuga - Nyabu	6	6	-1.99	30.193	Masaka Rd_1	0.5573654	5.254886994	1.33767696	Start	Very Short	Close	normal
8	Kabuga	Kabuga - Nyabu	7	7	-1.985	30.191	At 19_Inyangi	1.173229832	5.812252394	2.815751596	Start	Short	Medium	normal
9	Kabuga	Kabuga - Nyabu	8	8	-1.979	30.182	KK 3 Rd_15	0.219123292	6.985482226	0.525895901	Early	Very Short	Very Close	normal
10	Kabuga	Kabuga - Nyabu	9	9	-1.977	30.182	Kibaya Minagr	0.534957457	7.204605518	1.283897896	Early	Very Short	Close	normal
11	Kabuga	Kabuga - Nyabu	10	10	-1.974	30.179	Bus Route 107	0.413880807	7.739562975	0.993313937	Early	Very Short	Very Close	normal
12	Kabuga	Kabuga - Nyabu	11	11	-1.971	30.176	Mulindi1	0.714276414	8.153443782	1.714263393	Early	Very Short	Close	normal
13	Kabuga	Kabuga - Nyabu	12	12	-1.969	30.17	KK 3 Rd_11	0.754819859	8.867720195	1.811567663	Early	Very Short	Close	normal
14	Kabuga	Kabuga - Nyabu	13	13	-1.967	30.163	At 15 Ndera1	1.774549572	9.622540055	4.258918973	Early	Short	Medium	normal
15	Kabuga	Kabuga - Nyabu	14	14	-1.958	30.15	SEZ B2	1.309339478	11.39708963	3.142414747	Early	Short	Medium	normal
16	Kabuga	Kabuga - Nyabu	15	15	-1.954	30.139	Kigali Parents	0.852447952	12.7064291	2.045875085	Early	Short	Close	normal
17	Kabuga	Kabuga - Nyabu	16	16	-1.957	30.132	La Parisse1	1.017854079	13.55887706	2.442849791	Mid	Short	Medium	normal
18	Kabuga	Kabuga - Nyabu	17	17	-1.963	30.125	Ku Cyamutzig	1.070272592	14.57673114	2.568654221	Mid	Short	Medium	normal
19	Kabuga	Kabuga - Nyabu	18	18	-1.96	30.115	Prince House2	0.372548431	15.64700373	0.894116234	Mid	Very Short	Very Close	normal
20	Kabuga	Kabuga - Nyabu	19	19	-1.961	30.112	Alpha Palace F	0.396151284	16.01955216	0.950763083	Mid	Very Short	Very Close	normal
21	Kabuga	Kabuga - Nyabu	20	20	-1.963	30.109	Good Year2	1.006840164	16.41570344	2.416416393	Mid	Short	Medium	normal
22	Kabuga	Kabuga - Nyabu	21	21	-1.968	30.102	Sonatubes	0.419780487	17.42254361	1.007473168	Mid	Very Short	Very Close	normal
23	Kabuga	Kabuga - Nyabu	22	22	-1.969	30.098	Bralirwa3	0.42490554	17.84232409	1.019773296	Mid	Very Short	Very Close	normal
24	Kabuga	Kabuga - Nyabu	23	23	-1.968	30.094	Amasezerano	0.842559119	18.26722963	2.022141885	Mid	Short	Close	normal
25	Kabuga	Kabuga - Nyabu	24	24	-1.969	30.087	Rwandex	0.472096312	19.10978875	1.133031148	Mid	Very Short	Very Close	normal
26	Kabuga	Kabuga - Nyabu	25	25	-1.968	30.083	Kwa Mironko2	0.400322989	19.58188507	0.960775173	Late	Very Short	Very Close	normal
27	Kabuga	Kabuga - Nyabu	26	26	-1.965	30.08	KN 3 Rd_6	0.408243289	19.98220805	0.979783893	Late	Very Short	Very Close	normal
28	Kabuga	Kabuga - Nyabu	27	27	-1.962	30.078	Volta Super1	0.601627875	20.39045134	1.443906901	Late	Very Short	Close	normal
29	Kabuga	Kabuga - Nyabu	28	28	-1.957	30.076	Ku Maz1	0.794848627	20.99207922	1.907636704	Late	Very Short	Close	normal
30	Kabuga	Kabuga - Nyabu	29	29	-1.95	30.075	Chez Rasta	0.415276558	21.78692784	0.99666374	Late	Very Short	Very Close	normal
31	Kabuga	Kabuga - Nyabu	30	30	-1.947	30.074	One Love	0.801787986	22.2022044	1.924291166	Late	Very Short	Close	normal
32	Kabuga	Kabuga - Nyabu	31	31	-1.941	30.069	Forex Form_P	0.723142135	23.00399239	1.735541123	Late	Very Short	Close	normal
33	Kabuga	Kabuga - Nyabu	32	32	-1.937	30.064	Kigali Gaz_Kin	0.45730886	23.72713452	1.097541265	Late	Very Short	Very Close	normal

CHALLENGES & LESSONS LEARNED



Challenges Encountered:

- Interpreting mixed coordinate formats.
- Standardizing stop names without official references.
- Handling incomplete route information.
- Ensuring spatial accuracy for distance calculations



Lessons Learned:

- Data preprocessing is crucial (80% of data science work)
- Clean, structured data enables effective CI applications
- Domain knowledge essential for data interpretation
- Quality preprocessing directly impacts model performance

Conclusion: The preprocessed dataset is now ready for: Machine learning model training, Route optimization algorithms, Real-time arrival prediction systems, Intelligent transportation analytics.

TRANSCONNECT DATA PREPROCESSING SUMMARY

====

Original stops: 36

Cleaned stops: 36

Total route distance: 26.07 km

Total travel time: 62.6 minutes

Augmented dataset size: 216 rows

Conditions simulated: 6

Summary for Transconnect for Route "**Kabuga – Nyabugogo**"

CODE DEMONSTRATION STRUCTURE

Python Script Sections:

- Data Loading & Initial Assessment
- Data Cleaning Operations
- Data Integration & Deduplication
- Feature Engineering (Distance Calculations)
- Data Transformation & Discretization
- Data Augmentation for Model Training
- Visualization of Results

Tools & Libraries:

- Pandas for data manipulation
- NumPy for numerical operations
- Scikit-learn for preprocessing
- Haversine for distance calculations
- Matplotlib/Plotly for visualization
- Folium for map-based displays

Codes:

<https://drive.google.com/file/d/120uvRQOfj2I6KP-9buP23hwFub0JSyPO/view?usp=sharing>

Dataset:

https://docs.google.com/spreadsheets/d/1nQS_Vadk_r9I_5k6FVnBV3_UW04ITMII-7orBFDDIBv8/edit?usp=sharing

AI MODELS

Selected Models for TransConnect



1. Linear Regression – simple baseline model



2. Random Forest – stronger non-linear model



Linear Regression – What Was Completed

Provides a simple baseline for ETA prediction

Key Features:

-  Simple, Fast, and Computationally Efficient
-  Highly Interpretable
-  Works Well on Linearly Related Data (travel time, distance, time of day, road features)



Advantages: easy, fast, interpretable



Disadvantages: weak with non-linear route patterns

Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
from sklearn.preprocessing import StandardScaler
```

Set Model

```
def __init__(self):
    self.model = LinearRegression()
    self.results = {}
    self.scaler = StandardScaler()
```

Train Dataset

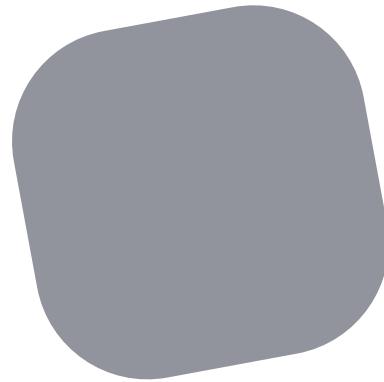
```
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.2, random_state=42)
X_train_scaled = self.scaler.fit_transform(X_train)
X_test_scaled = self.scaler.transform(X_test)

self.model.fit(X_train_scaled, y_train)
y_pred = self.model.predict(X_test_scaled)
```

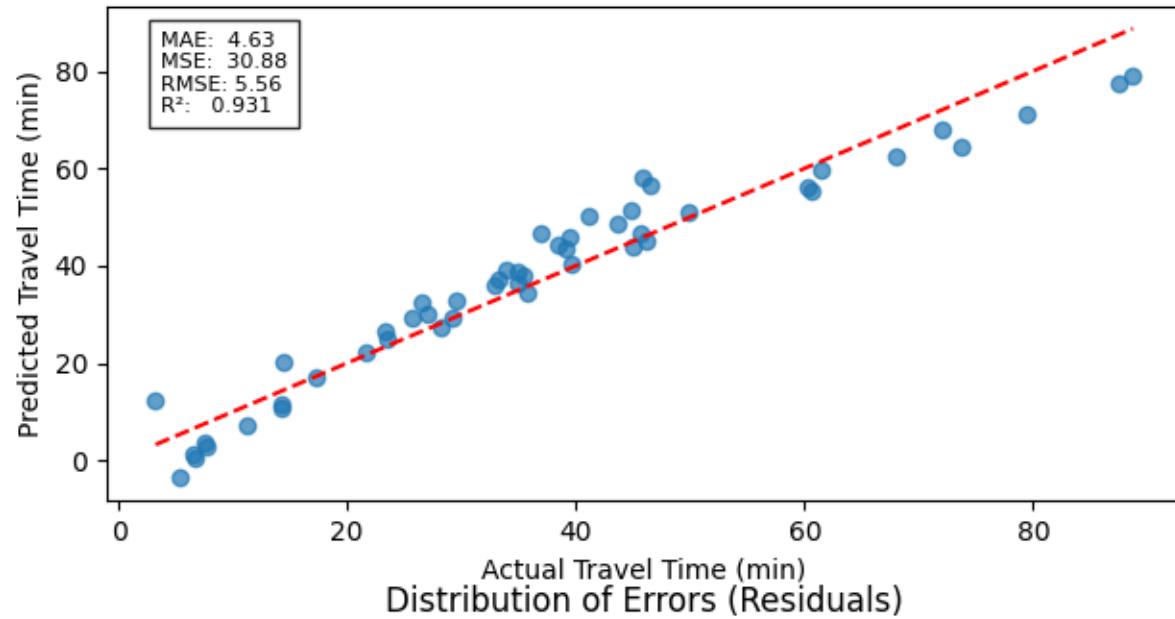
Performance Metric

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```


VISUALIZATION




Actual vs Predicted (Linear Regression)



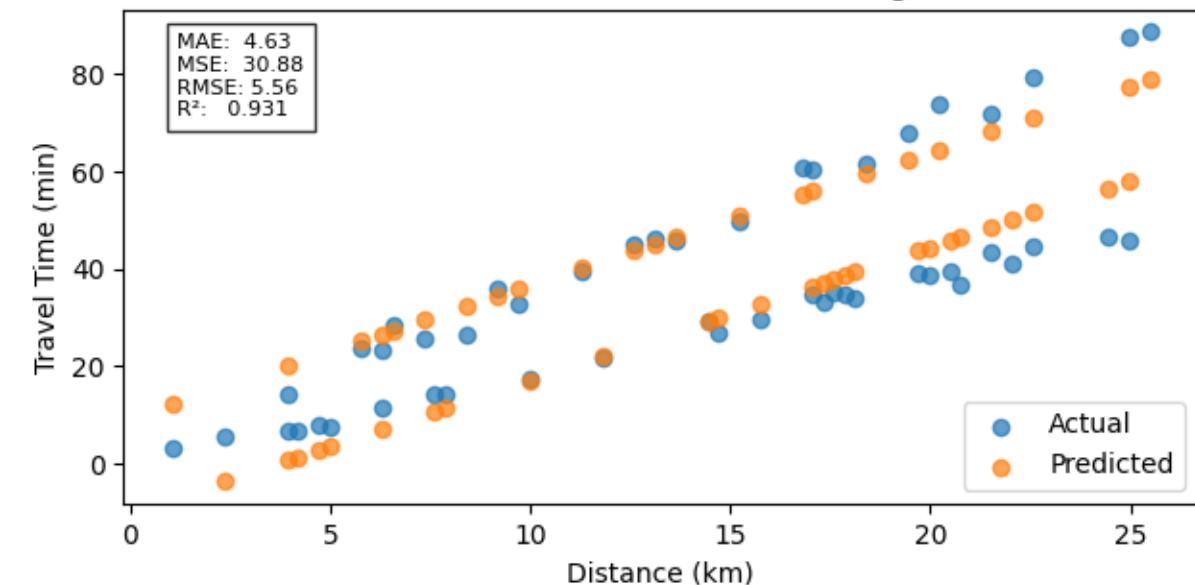
Residual Plot



Distribution of Errors (Residuals)



Predicted vs Actual Travel Time Along Route



Random Forest – What Was Completed

Handles non-linear transit patterns well

Key Features:

- Features do not need normalization or standardization (unlike Linear Regression)
- Provides Feature Importance
- High Accuracy on Tabular Data

- **Advantages:** high accuracy, robust
- **Disadvantages:** slower, harder to interpret

CODE SNIPPET

Library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

Model Fitting

```
def __init__(self):
    self.model = RandomForestRegressor(
        n_estimators=150,
        random_state=42,
        n_jobs=-1
    )
    self.results = {}
```

Train Model

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

self.model.fit(X_train, y_train)
y_pred = self.model.predict(X_test)
```


Performance Metric

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```


VISUALIZATION




Actual vs Predicted (Random Forest)



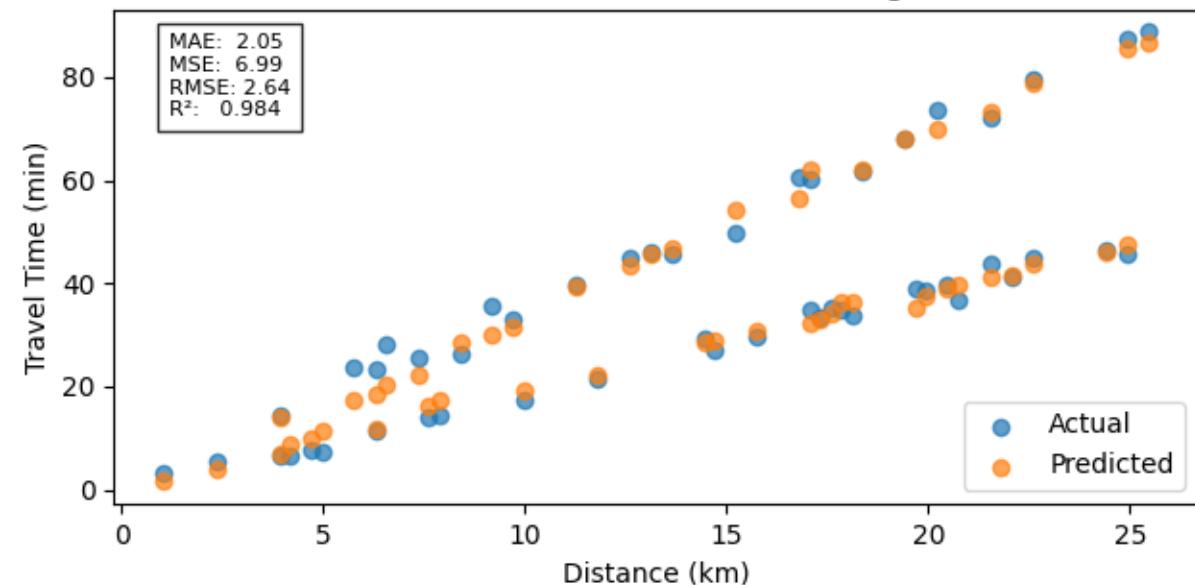
Distribution of Errors (Residuals)



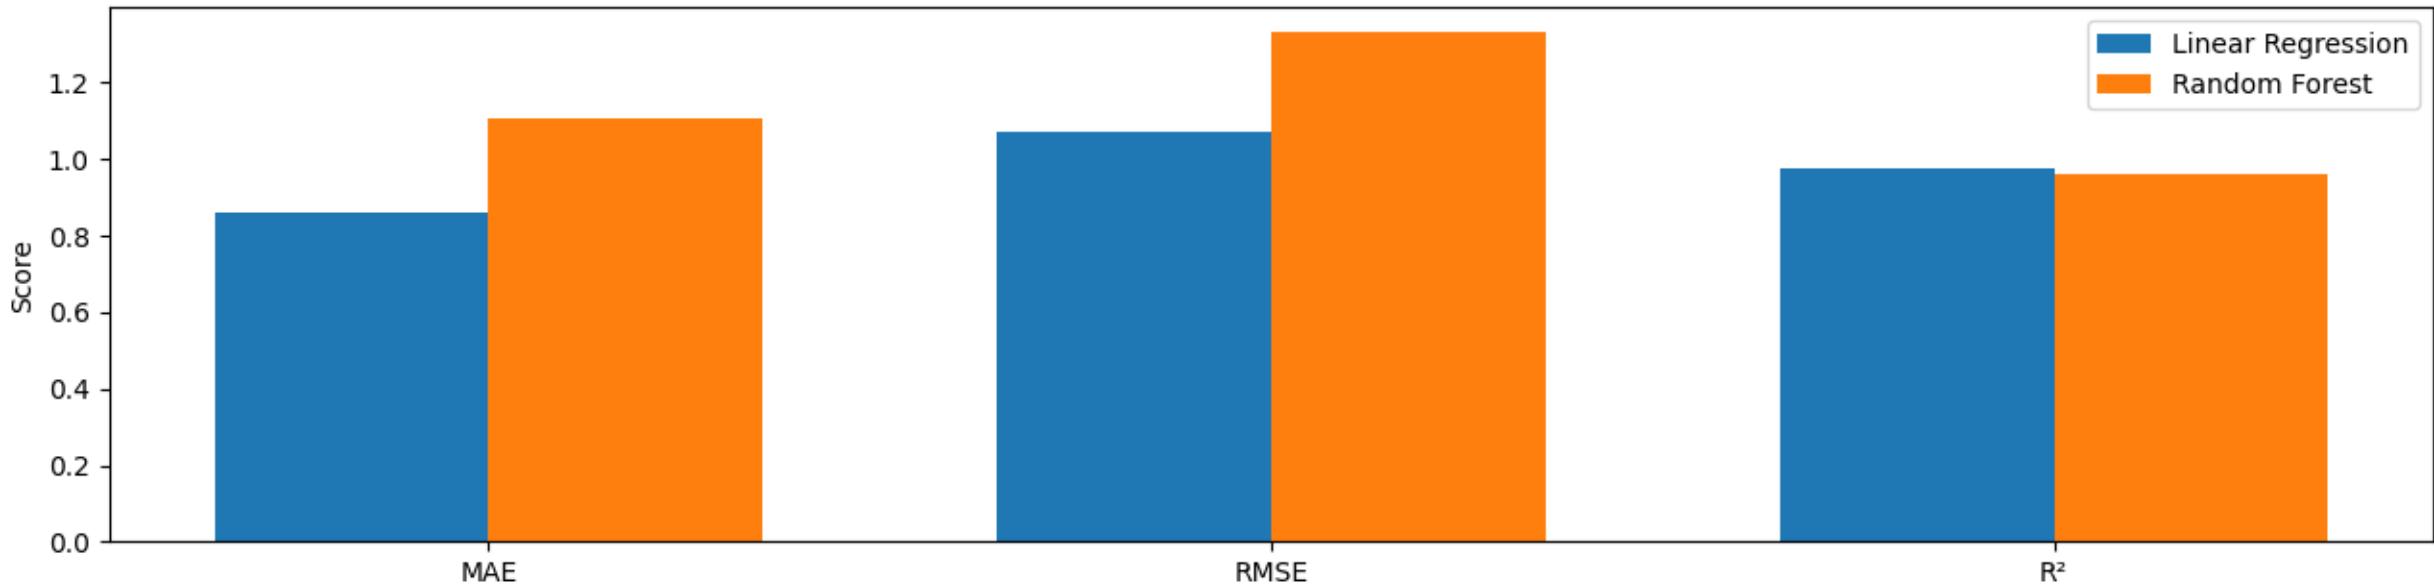
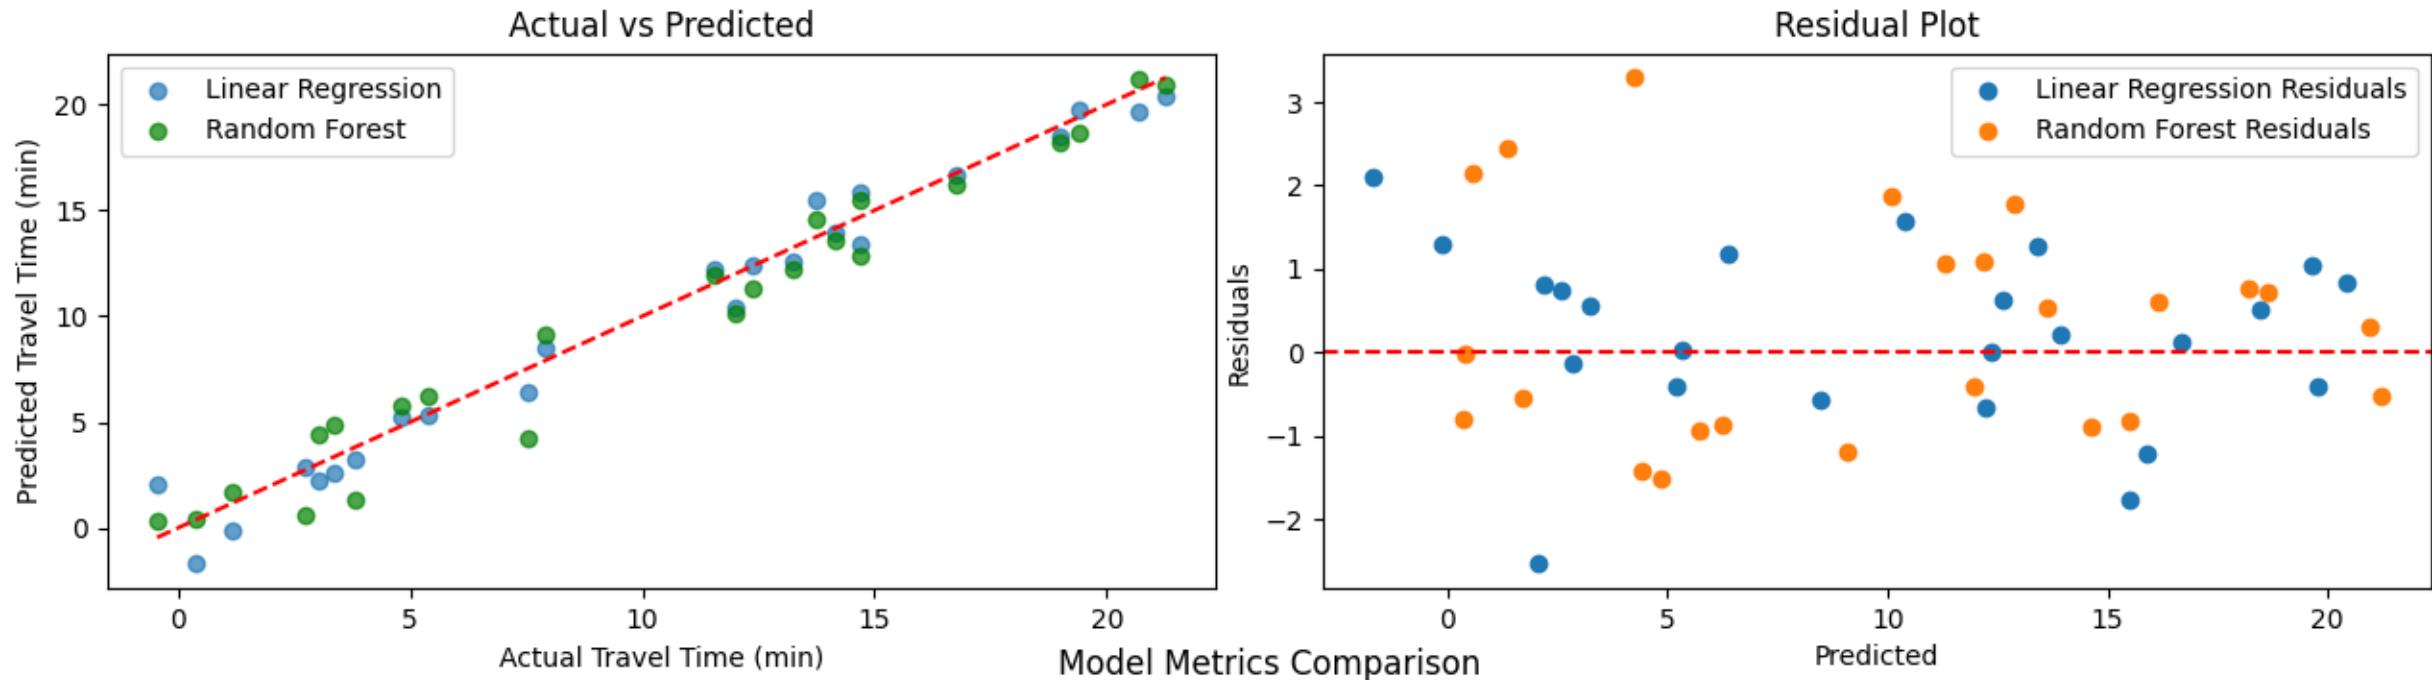
Residual Plot (Random Forest)



Predicted vs Actual Travel Time Along Route



LINEAR REGRESSION VS RANDOM FOREST



Cont...

Selected Models for TransConnect



3. Gradient Boosting Regression – Ensemble model of many decision trees



4. Decision Tree – supervised non-linear model

CASE STUDY

ROUTE

Kabuga - Nyabugogo



Gradient Boosting – What Was Completed

- Learns sequentially: each tree corrects previous errors
- Very powerful for structured data

Key Features:

-  High-accuracy ETA prediction
-  Modeling traffic effects on travel time
-  Works Well on Linearly Related Data (travel time, distance, time of day, road features)



Advantages:

- High prediction accuracy
- Handles complex relationships
- Great for time-based forecasting



Disadvantages:

- Slower training
- Sensitive to hyperparameters
- Can overfit if not tuned

Code Snippet

Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score
```

Set Model

```
model = GradientBoostingRegressor(
    n_estimators=300,
    learning_rate=0.05,
    max_depth=4,
    random_state=42
)
```

Train Dataset

```
X = df[["distance", "is_traffic"]]
y = df["travel_time"]


X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state=42)

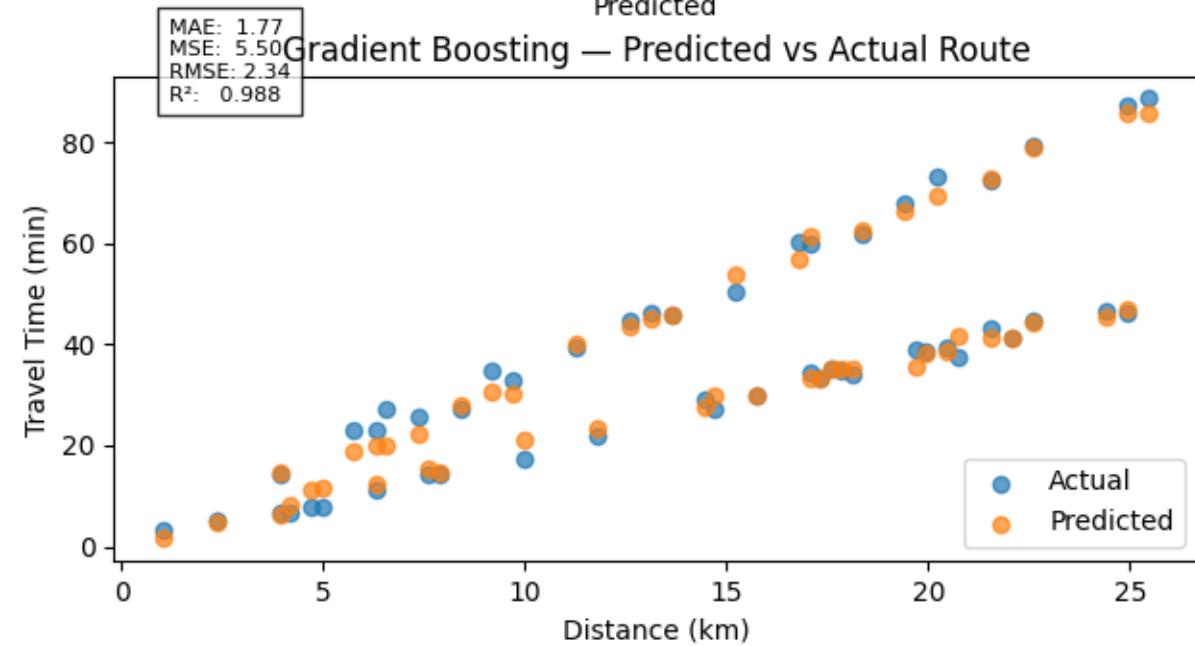
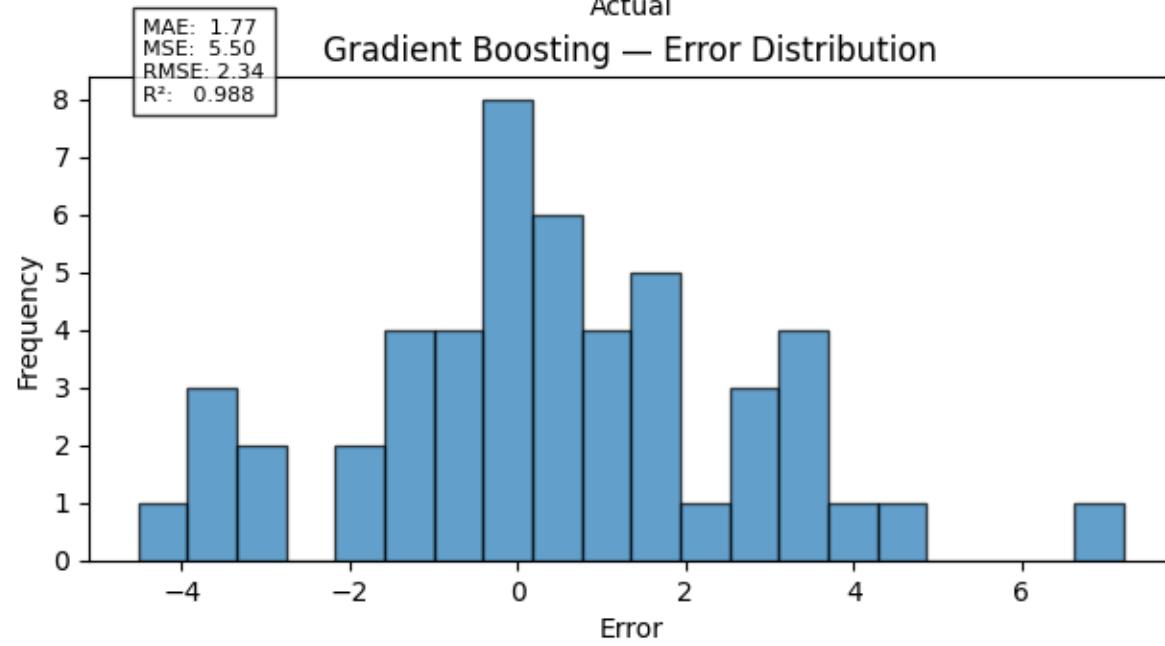
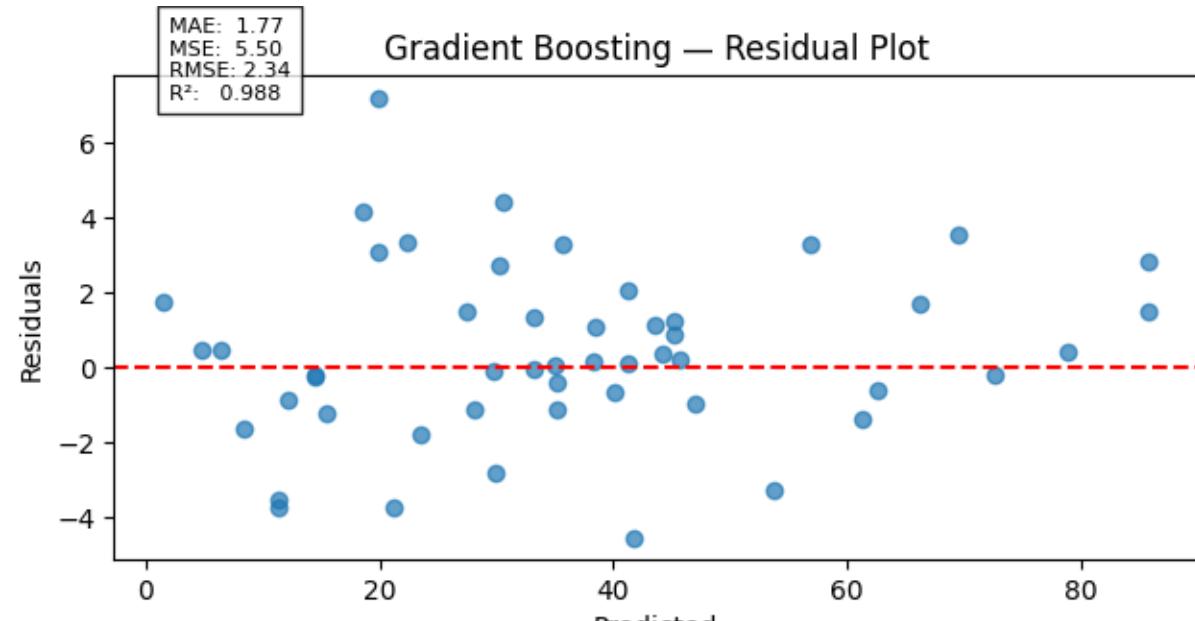
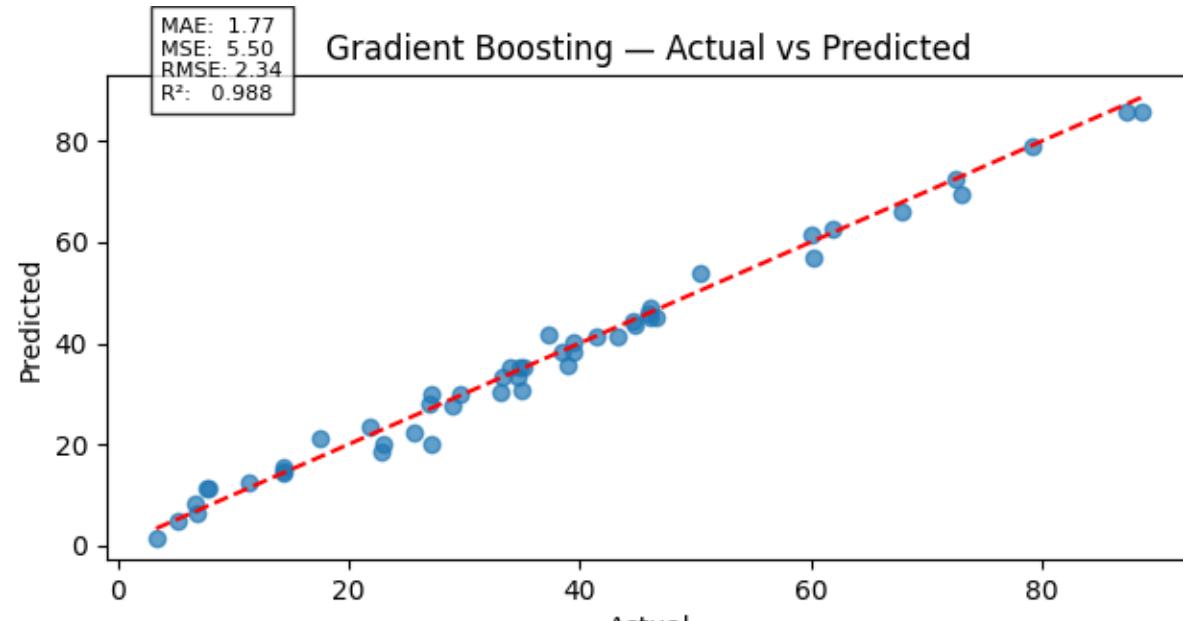
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
residuals = y_test - y_pred
```

Performance Metric

```
mae   = mean_absolute_error(y_test, y_pred)
mse   = mean_squared_error(y_test, y_pred)
rmse  = np.sqrt(mse)
r2    = r2_score(y_test, y_pred)
```

VISUALIZATION





Decision Tree – What Was Completed

- Supervised non-linear model
- Splits data into rules based on thresholds
- Easy to interpret and visualize

Key Features:

- Predicting travel time by condition (traffic vs normal)
- Detecting unusual stops or delays
- Fast to train, good for quick predictions

- **Advantages:**

- Simple and interpretable
- Handles non-linear data
- Works well with small datasets

- **Disadvantages:**

- Can overfit without pruning
- Unstable with small changes in data

CODE SNIPPET

Library

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error,
mean_squared_error, r2_score
```

Model Fitting

```
model = DecisionTreeRegressor(max_depth=6,
                               random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
residuals = y_test - y_pred
```

Train Model

```
X = df[["distance", "is_traffic"]]
y = df["travel_time"]


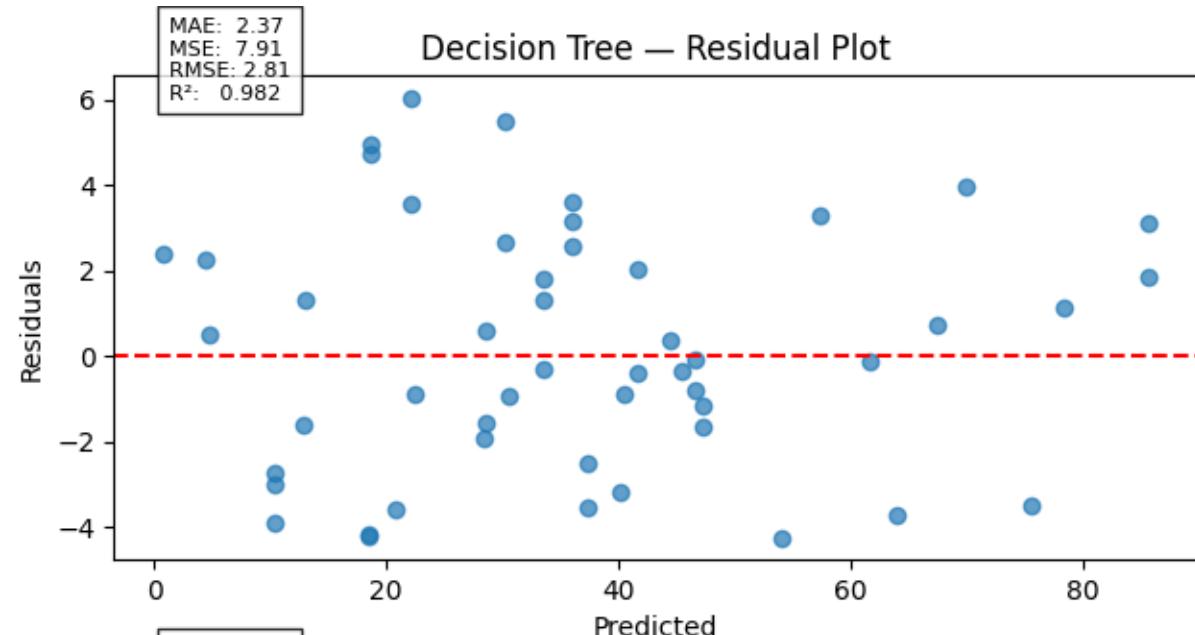
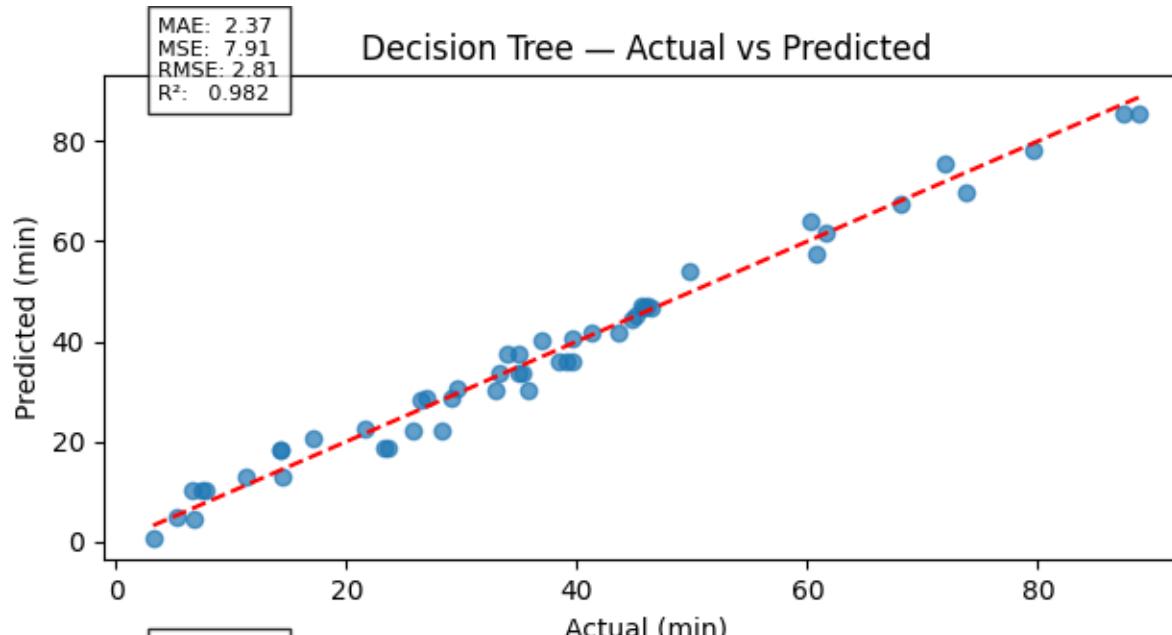
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25, random_state=42)
```

Performance Metric

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```

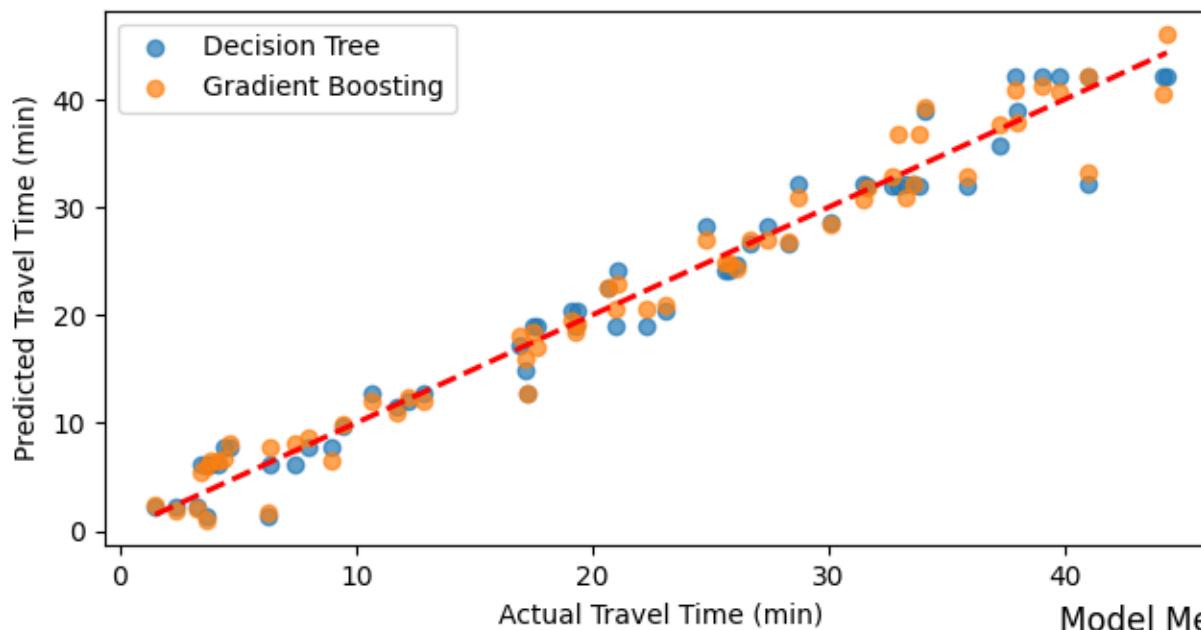
VISUALIZATION






GRADIENT BOOSTING REGRESSION VS DECISION TREE


Actual vs Predicted (DT vs GB)



Residuals Comparison



Model Metrics Comparison



SUMMARY

Decision Tree & Gradient Boosting

- Good for quick rule-based predictions
- Easy to interpret
- Best accuracy
- Ideal for real-time ETA predictions

RNN Model for TransConnect



- 5. **LSTM** (Long Short-Term Memory)
deep learning model for sequential data



Why LSTM?

- Designed for sequential data
- Learns traffic patterns over time
- Captures rush hour vs normal variations
- More accurate than standard ML models

■ Model Structure:

- **Input** (distance, is_traffic, hour)
- LSTM **Layer** (64 units) – learns sequential behavior
- **Dropout** (0.2) – prevents overfitting
- Dense Layer (32 neurons)
- Output Layer: 1 node (Predicted travel time)


Key Features:

- Handles time-dependent patterns
- Smooth prediction curve instead of noisy output
- Works well with irregular travel speeds
- Adjust predictions using context (distance, traffic, hour)

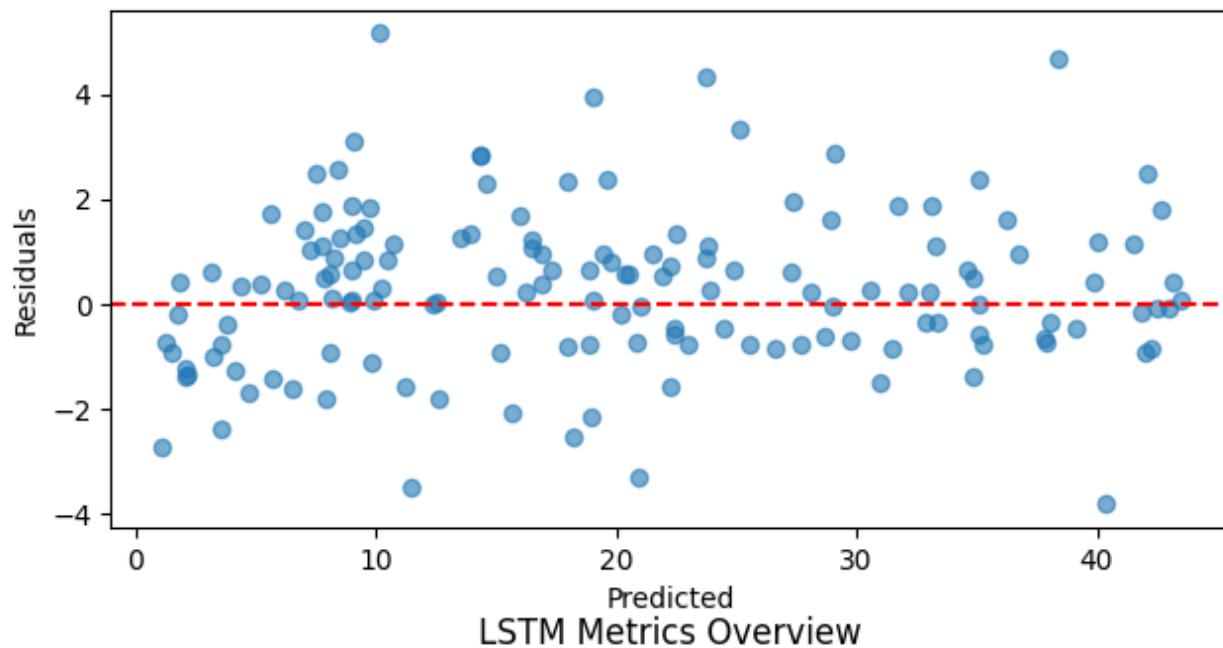
VISUALIZATION




Actual vs Predicted (LSTM)



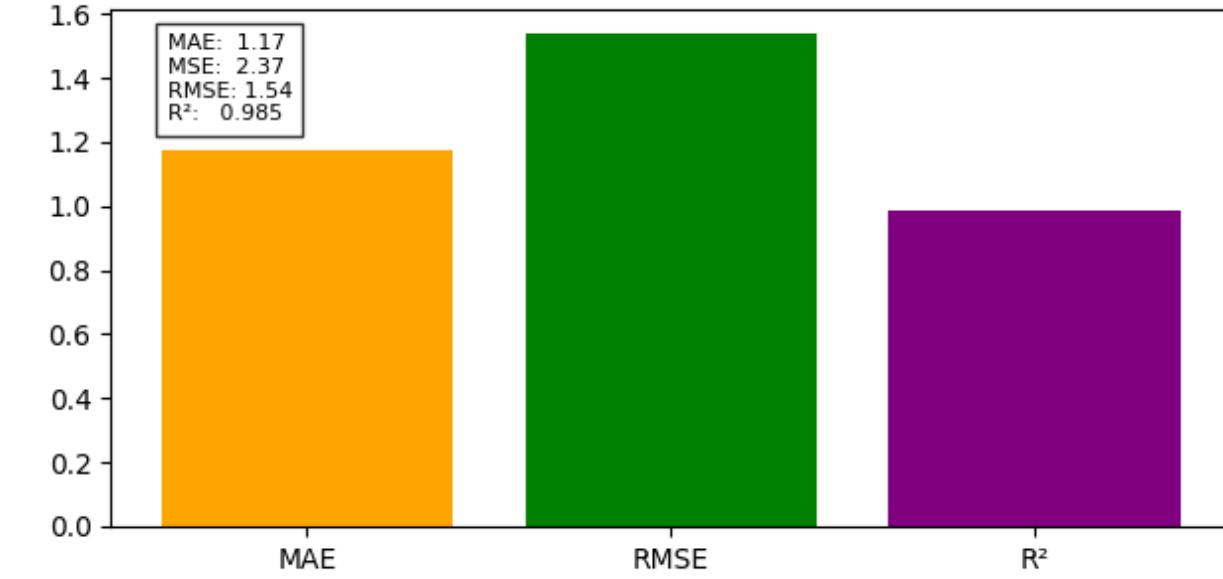
Residual Plot (LSTM)



LSTM Training Curve



LSTM Metrics Overview





BEST MODEL – AMONG ALL

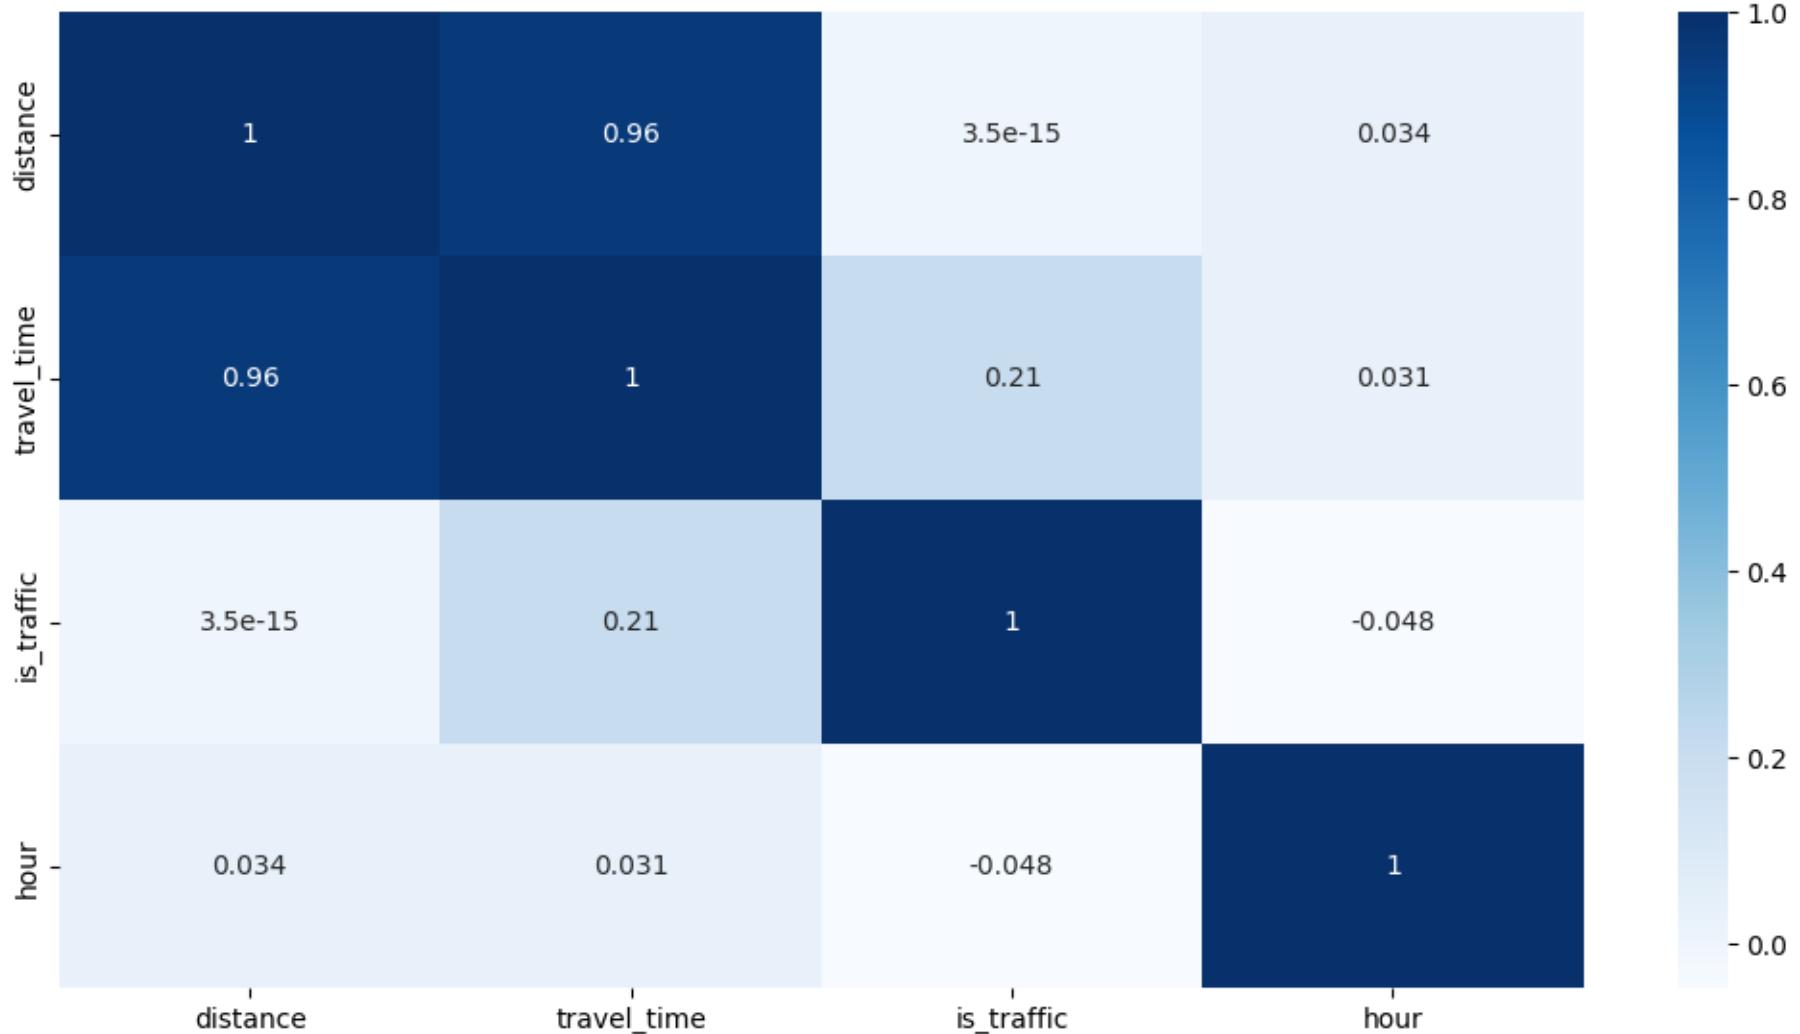
LSTM

`best_transconnect_model.h5`

MORE COMPARISON

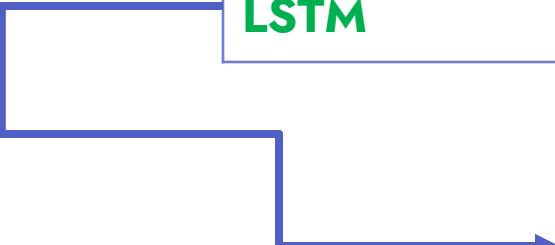


TransConnect Feature Correlation Matrix



Model Performance Summary

Model	MAE	MSE	RMSE	R ²
Linear Regression	1.6956	4.5832	2.1408	0.97095
Decision Tree	1.5216	4.0942	2.0234	0.97405
Random Forest	1.4514	3.5948	1.8960	0.97721
Gradient Boosting	1.4090	3.2546	1.8040	0.97937
LSTM	1.2463	2.7779	1.6667	0.98239

- 
- 
- Highest R²
 - lowest RMSE
 - learns sequential traffic patterns

Challenges

- Limited historical data
- GPS noise and inconsistent-format coordinates
- Few features available for prediction
- Hard to validate ETA accuracy in real time


Mitigation Measures



Collect more real bus timing logs (according to passengers, like us)



Apply smoothing & cleaning to GPS streams



Add engineered features (traffic, rush hour, holi



Test across multiple routes for reliability

END